

networkx: A High-Performance Graph Library for Python

Matthew Treinish¹, Ivan Carvalho², Georgios Tsilimigkounakis³, and Nahum Sá⁴

¹ IBM Quantum, IBM T.J. Watson Research Center, Yorktown Heights, USA

² University of British Columbia, Kelowna, Canada

³ National Technical University of Athens, Athens, Greece

⁴ Centro Brasileiro de Pesquisas Físicas, Rio de Janeiro, Brazil

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

In *networkx*, we provide a high-performance, flexible graph library for Python. *networkx* is inspired by *NetworkX* but addresses many performance concerns of the latter. *networkx* is written in Rust and is particularly suited for performance-sensitive applications that use graph representations.

Editor: [Vincent Knight](#) ↗

Reviewers:

- [@szhorvat](#)
- [@inakleinbottle](#)

Submitted: 28 October 2021

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Statement of need

networkx is a general-purpose graph theory library focused on performance. It wraps low-level Rust code ([Matsakis & Klock, 2014](#)) into a flexible Python API, providing fast implementations for graph data structures and popular graph algorithms.

networkx is inspired by the *NetworkX* library ([Hagberg et al., 2008](#)), but meets the needs of users that also need performance. Even though *NetworkX* is the de-facto standard graph and network analysis library for Python, it shows performance concerns. *NetworkX* prefers pure Python implementations, which leads to bottlenecks in computationally intensive applications that use graph algorithms.

networkx addresses those performance concerns by switching to a Rust implementation. It has support for shortest paths, isomorphism, matching, multithreading via rayon ([Stone & Matsakis, 2021](#)), and much more.

Related work

The graph and network analysis ecosystem for Python is rich, with many libraries available. *igraph* ([Csardi & Nepusz, 2006](#)), *graph-tool* ([Peixoto, 2014](#)), *SNAP* ([Leskovec & Sosič, 2016](#)), and *Networkit* ([Staudt et al., 2016](#)) are Python libraries written in C or C++ that can replace *NetworkX* with better performance. We also highlight *SageMath*'s graph theory module ([The Sage Developers, 2020](#)), which has a stronger focus in mathematics than *NetworkX*.

However, the authors found that no library matched the flexibility that *NetworkX* provided for interacting with graphs. *igraph* is efficient for static large graphs, but does not handle graph updates as efficiently. *SNAP* and *Networkit* do not support associated edge data with arbitrary Python types. *graph-tool* supports associated edge data at the cost of maintaining the data in a separate data structure. Thus, the main contribution of *networkx* is keeping the ease of use of *NetworkX* without sacrificing performance.

We note that existing code using *NetworkX* needs to be modified to use *networkx*. *networkx* is not a drop-in replacement for *NetworkX*, which may be a possible limitation for some users. The authors provide a *NetworkX* to *networkx* conversion guide on the documentation to aid in those situations.

Graph data structures

networkx provides two core data structures: `PyGraph` and `PyDiGraph`. They correspond to undirected and directed graphs, respectively. Graphs describe a set of nodes and the edges connecting pairs of those nodes. Internally, *networkx* leverages the *petgraph* library (bluss et al., 2021) to store the graphs using an adjacency list model and the *PyO3* library (Hewitt et al., 2021) for the Python bindings.

Nodes and edges of the graph may also be associated with data payloads. Payloads can contain arbitrary data, such as node labels or edge lengths. Common uses of data payloads include representing weighted graphs. Any Python object can be a data payload, which makes the library flexible because no assumptions are made about the payload types.

networkx operates on payloads with callbacks. Callbacks are functions that take payloads and return statically typed data. They resemble the named attributes in *NetworkX*. Callbacks are beneficial because they bridge the arbitrary stored data with the static types *networkx* expects.

A defining characteristic of *networkx* graphs is that each node maps to a non-negative integer node index, and similarly, each edge maps to an edge index. Those indices uniquely determine nodes and edges in the graph. Moreover, the indices provide a clear separation between the underlying graph structure and the data associated with payloads.




Figure 1: A Petersen graph, a hexagonal lattice graph, and a binomial tree graph created with `networkx.generators` and visualized with the `networkx.visualization` module.

Use Cases

networkx is suitable for modeling graphs ranging from a few nodes scaling up to millions. The library is particularly suited for applications that have core routines executing graph algorithms. In those applications, the performance of *networkx* considerably reduces computation time. Examples of applications using *networkx* include the Qiskit compiler (Treinish et al., 2021), PennyLane (Bergholm et al., 2020), atompack (Ullberg, 2021), and qtcodes (Jha et al., 2021).

For common use cases, *networkx* can provide speedups ranging from 3x to 100x faster compared to the same code using *NetworkX* while staying competitive with other compiled libraries like *igraph* and *graph-tool*. The gains in performance are application-specific, but as a general rule, the more work that is offloaded to *networkx* and Rust, the larger are the gains.

We illustrate use cases with examples from the field of quantum computing that motivated the development of the library.

Graph Creation, Manipulation, and Traversal

The first use case is based on the manipulation of directed acyclic graphs (DAGs) by Qiskit using *networkx*. Qiskit represents quantum circuits as DAGs which the compiler operates on to perform analysis and transformations (Childs et al., 2019).




Figure 2: Quantum circuit and its equivalent representation as a DAG of instructions built by Qiskit.

Qiskit creates a DAG whose nodes represent either instructions or registers present in the quantum circuit (Cross et al., 2021) and whose edges represent the registers each instruction operates on. Qiskit also applies transformations to the instructions, which manipulates the graph by adding and removing nodes and edges. *networkx* brings the graph data structure underlying those operations.

In addition, Qiskit needs to traverse the graph. Some transformations, such as greedily merging instructions to reduce circuit depth, require graph traversal. *networkx* offers the methods for traversals such as breadth-first search, depth-first search, and topological sorting.

Subgraph Isomorphism

The second use case is based on qubit mapping problem for Noisy Intermediate-Scale Quantum (NISQ) devices (Bharti et al., 2022; Preskill, 2018). NISQ devices do not have full connectivity among qubits, hence Qiskit needs to take into account an undirected graph representing the connectivity of the device when compiling quantum circuits. Qiskit transforms the quantum circuit such that the pairs of qubits executing two-qubit gates respect the device's architectural constraints. There are many proposed solutions to the qubit mapping problem, including algorithms based on subgraph isomorphism (Li et al., 2021).




Figure 3: Graph representing the connectivity of the **ibmq_montreal** device. Qiskit can check if the required connectivity by a circuit is subgraph isomorphic to the device's connectivity when solving the qubit mapping problem.

networkx implements the VF2 algorithm (Cordella et al., 2004) and some of the VF2++ heuristics (Jüttner & Madarasi, 2018) to solve subgraph isomorphism. The implementations include both checking if a mapping exists and returning a mapping among the nodes. Qiskit leverages *networkx* to provide an experimental layout method based on VF2. Qiskit checks if the graph representing the connectivity required by the circuit and the connectivity provided by

the device are subgraph isomorphic. If they are, Qiskit uses VF2 mapping to map the qubits without increasing the depth of the circuit.

Acknowledgements

We thank Kevin Krsulich for his help in getting *networkx* ready for use by Qiskit; Lauren Capelluto and Toshinari Itoko for their continued support and help with code review; and all of the *networkx* contributors who have helped the library improve over time.

References

- Bergholm, V., Izaac, J., Schuld, M., Gogolin, C., Alam, M. S., Ahmed, S., Arrazola, J. M., Blank, C., Delgado, A., Jahangiri, S., McKiernan, K., Meyer, J. J., Niu, Z., Száva, A., & Killoran, N. (2020). *PennyLane: Automatic differentiation of hybrid quantum-classical computations*. <https://arxiv.org/abs/1811.04968>
- Bharti, K., Cervera-Lierta, A., Kyaw, T. H., Haug, T., Alperin-Lea, S., Anand, A., Degroote, M., Heimonen, H., Kottmann, J. S., Menke, T., Mok, W.-K., Sim, S., Kwek, L.-C., & Aspuru-Guzik, A. (2022). Noisy intermediate-scale quantum algorithms. *Reviews of Modern Physics*, 94(1). <https://doi.org/10.1103/revmodphys.94.015004>
- bluss, Comets, J.-M., Borgna, A., Larralde, M., Mitchener, B., & Kochkov, A. (2021). Petgraph. In *GitHub repository*. GitHub. <https://github.com/petgraph/petgraph>
- Childs, A. M., Schoute, E., & Unsal, C. M. (2019). Circuit Transformations for Quantum Architectures. In W. van Dam & L. Mancinska (Eds.), *14th conference on the theory of quantum computation, communication and cryptography (TQC 2019)* (Vol. 135, pp. 3:1–3:24). Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik. <https://doi.org/10.4230/LIPIcs.TQC.2019.3>
- Cordella, L. P., Foggia, P., Sansone, C., & Vento, M. (2004). A (sub)graph isomorphism algorithm for matching large graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 26(10), 1367–1372. <https://doi.org/10.1109/TPAMI.2004.75>
- Cross, A. W., Javadi-Abhari, A., Alexander, T., Beaudrap, N. de, Bishop, L. S., Heidel, S., Ryan, C. A., Smolin, J., Gambetta, J. M., & Johnson, B. R. (2021). *OpenQASM 3: A broader and deeper quantum assembly language*. <https://doi.org/10.48550/arXiv.2104.14722>
- Csardi, G., & Nepusz, T. (2006). The igraph software package for complex network research. *InterJournal, Complex Systems*, 1695. <https://igraph.org>
- Hagberg, A. A., Schult, D. A., & Swart, P. J. (2008). Exploring network structure, dynamics, and function using NetworkX. In G. Varoquaux, T. Vaught, & J. Millman (Eds.), *Proceedings of the 7th python in science conference* (pp. 11–15). http://conference.scipy.org/proceedings/SciPy2008/paper_2/
- Hewitt, D., Kanagawa, Y., Kim, N., Grunwald, D., Niederbühl, A., messense, Kolenbrander, B., Brandl, G., & Ganssle, P. (2021). PyO3. In *GitHub repository*. GitHub. <https://github.com/PyO3/pyo3>
- Jha, S., Chen, J., Householder, A., & Mi, A. (2021). Qtcodes. In *GitHub repository*. GitHub. <https://github.com/yaleqc/qtcodes>
- Jüttner, A., & Madarasi, P. (2018). VF2++—an improved subgraph isomorphism algorithm. *Discrete Applied Mathematics*, 242, 69–81. <https://doi.org/10.1016/j.dam.2018.02.018>
- Leskovec, J., & Sosič, R. (2016). SNAP: A general-purpose network analysis and graph-mining library. *ACM Trans. Intell. Syst. Technol.*, 8(1). <https://doi.org/10.1145/2898361>

- 138 Li, S., Zhou, X., & Feng, Y. (2021). Qubit mapping based on subgraph isomorphism
139 and filtered depth-limited search. *IEEE Transactions on Computers*, 70(11), 1777–1788.
140 <https://doi.org/10.1109/tc.2020.3023247>
- 141 Matsakis, N. D., & Klock, F. S. (2014). The rust language. *Proceedings of the 2014*
142 *ACM SIGAda Annual Conference on High Integrity Language Technology*, 103–104. <https://doi.org/10.1145/2663171.2663188>
143
- 144 Peixoto, T. P. (2014). The graph-tool python library. *Figshare*. [https://doi.org/10.6084/m9.](https://doi.org/10.6084/m9.figshare.1164194)
145 [figshare.1164194](https://doi.org/10.6084/m9.figshare.1164194)
- 146 Preskill, J. (2018). Quantum computing in the NISQ era and beyond. *Quantum*, 2, 79.
147 <https://doi.org/10.22331/q-2018-08-06-79>
- 148 Staudt, C. L., Sazonovs, A., & Meyerhenke, H. (2016). NetworkKit: A tool suite for large-scale
149 complex network analysis. *Network Science*, 4(4), 508–530. [https://doi.org/10.1017/nws.](https://doi.org/10.1017/nws.2016.20)
150 [2016.20](https://doi.org/10.1017/nws.2016.20)
- 151 Stone, J., & Matsakis, N. D. (2021). Rayon: A data parallelism library for rust. In *GitHub*
152 *repository*. GitHub. <https://github.com/rayon-rs/rayon>
- 153 The Sage Developers. (2020). *SageMath, the Sage Mathematics Software System (Version*
154 *9.1)*. <https://doi.org/10.5281/zenodo.4066866>
- 155 Treinish, M., Gambetta, J., Rodríguez, D. M., Marques, M., Bello, L., Wood, C. J., Gomez,
156 J., Nation, P., Chen, R., Winston, E., Gacon, J., Cross, A., Krsulich, K., Sertage, I. F.,
157 Wood, S., Alexander, T., Capelluto, L., Puente González, S. de la, Rubio, J., ... Woerner,
158 S. (2021). *Qiskit/qiskit-terra: Qiskit terra 0.19.1* (Version 0.19.1) [Computer software].
159 Zenodo. <https://doi.org/10.5281/zenodo.2583252>
- 160 Ullberg, S. (2021). Atompac: A flexible python library for atomic structure generation. In
161 *GitHub repository*. GitHub. <https://github.com/seatonullberg/atompac>